

Big Data Engineering with Distributed Systems

Data Science Dojo

Agenda

- Introduction:
 - Data engineering for data scientists
 - The “5 Vs” of Big Data
- A key problem – machine learning at scale
- Distributed computing with Apache Hadoop & Hive
- Hadoop in the Azure cloud
- Machine learning at scale with Apache Mahout
- Distributed computing v2.0 – Apache Spark

Data Engineering for Data Scientists



Driving a car

VS



Servicing a car

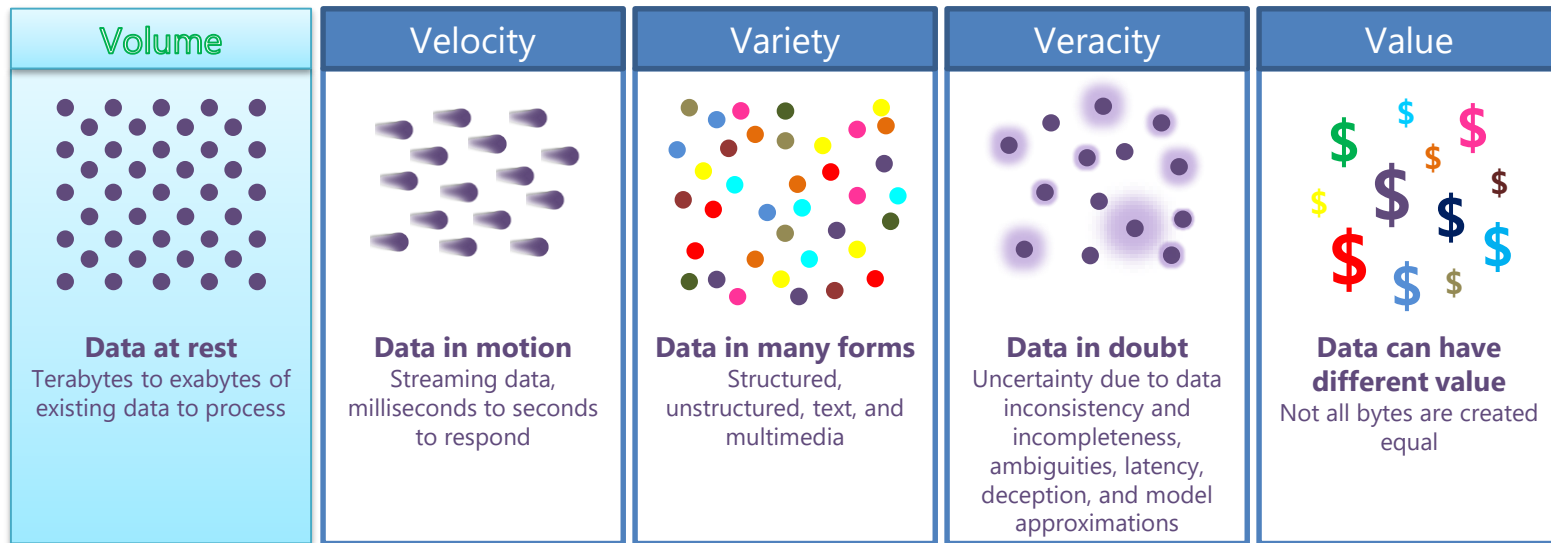
Goals:

- Teach you about data engineering topics/concepts

Non goals:

- Managing or administering a Hadoop cluster

5 Vs of Big Data





- **Goal:** As data scientists we want cost-effective access to the raw materials for our data products!

MACHINE LEARNING AT SCALE



OSS R Limits

- Single core
- Single threaded



Quad Core Laptop




Model A



Model B



Model C



OSS R Limits


- Single core
- Single threaded
- All in memory (RAM)
- Vectors & Matrices capped at 4,294,967,295 elements (rows) if 32-bit version; $2^{32} - 1$

OSS R Limits: RAM

- All in memory (RAM)

Max Data Limit = (Total RAM Access x 80%) – Normal RAM Usage

Laptop Example:



Max Data Limit = (5.9 gb x 80%) – 3.2gb
Max Data Limit = ~1.52gb

*R data frames actually bloats data files by 3x
R Data Limit = ~1.52gb ÷ 3 = ~506.7mb

OSS R Limits: RAM

INSTANCE	CORES	RAM	DISK SIZES ¹	PRICE
M64MS	64	1,750.00 GiB	2,000 GB	\$10.34/hr
M128S	128	2,000.00 GiB	4,000 GB	\$13.34/hr

Azure's VM with largest RAM*:

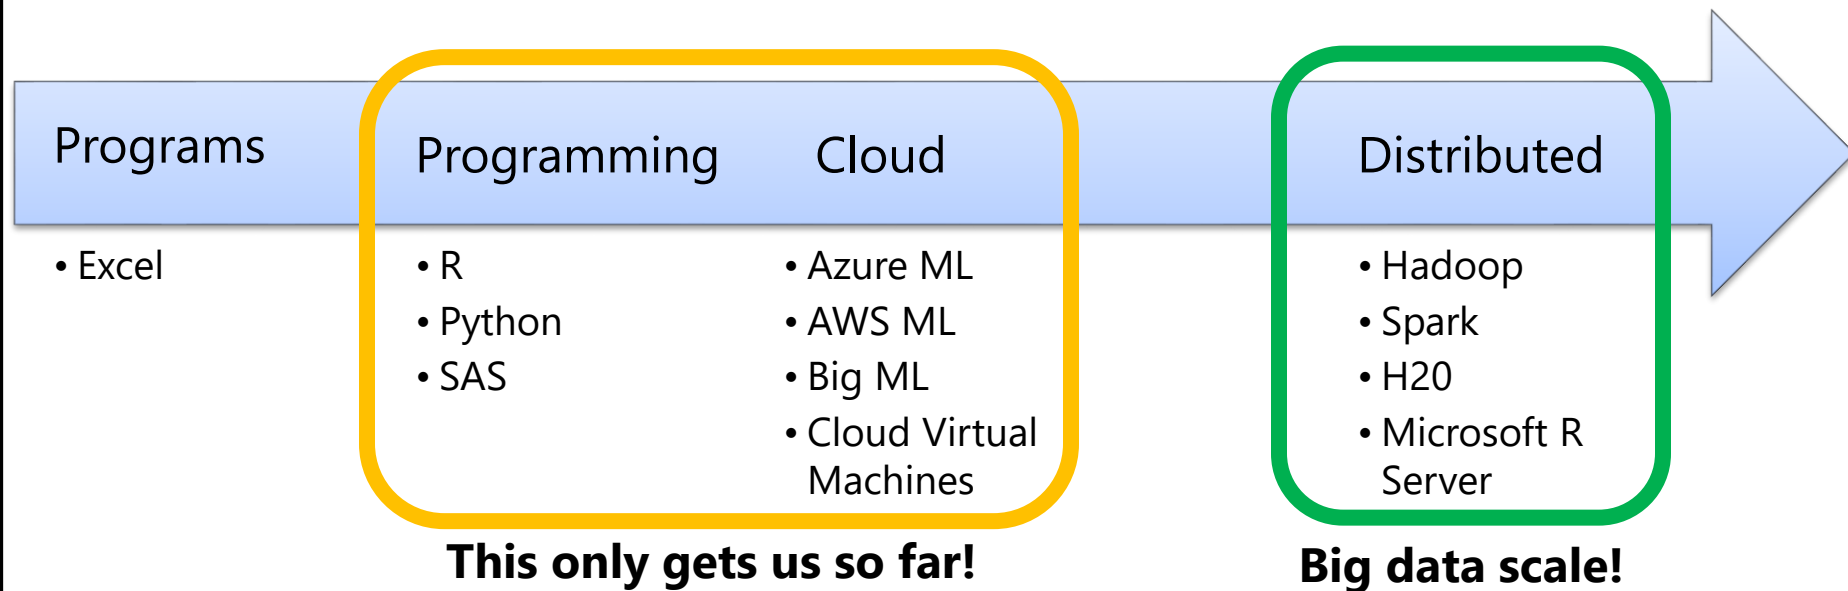
Max Data Limit = (2000gb x 80%) - 1gb

Max Data Limit = ~1600gb

R Data Limit = ~1600gb ÷ 3 = ~533.33 gb

24x7x52 Annual Cost: \$116,938.44!

Machine Learning Scaling





DISTRIBUTED COMPUTING WITH APACHE HADOOP

Turn Back The Clock, The Mainframe




- “Big Iron”
- Backbone of computing for decades.
- Still widely used.
- “Scale-up” model of shared computing.
- Core platform is cost effective, ecosystem is not (e.g., software licensing).
- The original VM host!

Distributed Computing



Cloud Computing




- Conceptually – a combination of mainframe and distributed computing.
- VM hosts are now the “Big Iron”.
- Many VMs work together to distribute workloads.
- Some workloads on dedicated HW (e.g., SAP HANA).

Scaling Computational Power



Old Scaling:

- Vertical Scaling, Scaling UP
- High performance computers




New Scaling:

- Horizontal Scaling, Scaling OUT
- Commodity hardware, distributed


What is Hadoop?

- OSS Platform for distributed computing over Internet-scale data.
- Originally built at Yahoo!
- Implementation of ideas (e.g., MapReduce) published by Google.
- The de facto standard big data platform.
- Named after a stuffed animal belonging to Doug Cutting's son.

hadoop




Hadoop at Base




Distributed batch processing engine for big data.

HDFS & MapReduce



Processing: 30 hours


HDFS & MapReduce



2 Computers

Processing: 15 hours


HDFS & MapReduce



Most Cases, Linear Scaling Of Processing Power


Number of Computers	Processing Time (hours)
1	30
2	15
3	10
4	7.5
5	6
6	5
7	4.26
8	3.75
9	3.33

If dogs were servers...






HDFS

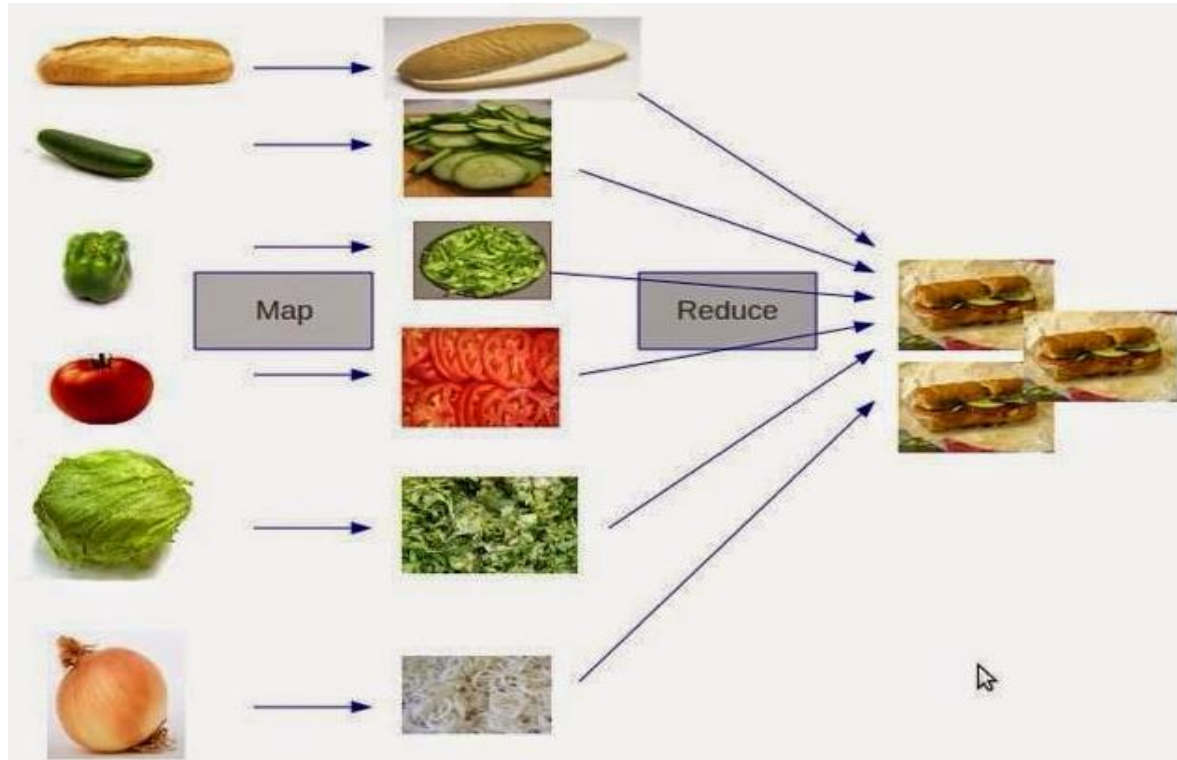
HDFS Partitioning



HDFS Redundancy



MapReduce – Sandwich Analogy



Limitations with MapReduce

- ~70 lines of code to do anything
- Slow
- Troubleshooting multiple computers
- Good devs are scarce
- Expensive certifications

```
1 package org.apache.hadoop.examples;
2
3 import java.io.IOException;
4 import java.util.StringTokenizer;
5
6 import org.apache.hadoop.conf.Configuration;
7 import org.apache.hadoop.fs.Path;
8 import org.apache.hadoop.io.IntWritable;
9 import org.apache.hadoop.io.Text;
10 import org.apache.hadoop.mapreduce.Job;
11 import org.apache.hadoop.mapreduce.Mapper;
12 import org.apache.hadoop.mapreduce.Reducer;
13 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
14 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
15 import org.apache.hadoop.util.GenericOptionsParser;
16
17 public class WordCount {
18
19     public static class TokenizerMapper
20         extends Mapper<Object, Text, Text, IntWritable>{
21
22         private final static IntWritable one = new IntWritable(1);
23         private Text word = new Text();
24
25         public void map(Object key, Text value, Context context
26             ) throws IOException, InterruptedException {
27             StringTokenizer itr = new StringTokenizer(value.toString());
28             while (itr.hasMoreTokens()) {
29                 word.set(itr.nextToken());
30                 context.write(word, one);
31             }
32         }
33     }
```

DISTRIBUTED COMPUTING WITH APACHE HIVE

What is Hive?

- Abstraction built on top of MapReduce & HDFS.
- Makes Hadoop look like an RDBMS (e.g., coding in SQL).
- Developed by Facebook to democratize Hadoop
- Applies structure to data at runtime ("schema on read").



Schema on Read



Data File



Unstructured
Data



Data File

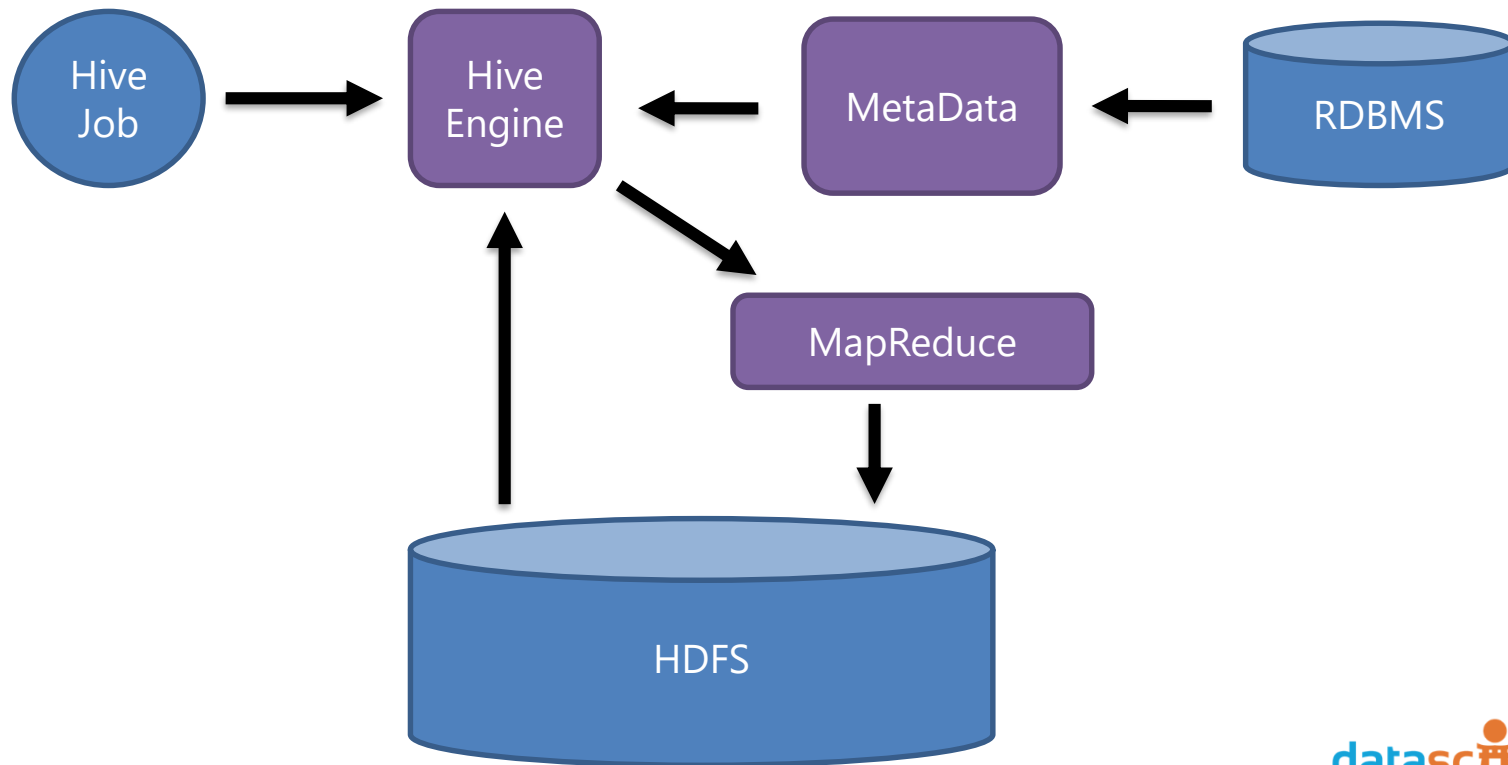


Metadata File/DB



Structured
Data

Hive Architecture



Hive Jobs



Word Count Revisited

```
1 package org.apache.hadoop.examples;
2
3 import java.io.IOException;
4 import java.util.StringTokenizer;
5
6 import org.apache.hadoop.conf.Configuration;
7 import org.apache.hadoop.fs.Path;
8 import org.apache.hadoop.io.IntWritable;
9 import org.apache.hadoop.io.Text;
10 import org.apache.hadoop.mapreduce.Job;
11 import org.apache.hadoop.mapreduce.Mapper;
12 import org.apache.hadoop.mapreduce.Reducer;
13 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
14 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
15 import org.apache.hadoop.util.GenericOptionsParser;
16
17 public class WordCount {
18
19     public static class TokenizerMapper
20         extends Mapper<Object, Text, Text, IntWritable>{
21
22         private final static IntWritable one = new IntWritable(1);
23         private Text word = new Text();
24
25         public void map(Object key, Text value, Context context
26             ) throws IOException, InterruptedException {
27             StringTokenizer itr = new StringTokenizer(value.toString());
28             while (itr.hasMoreTokens()) {
29                 word.set(itr.nextToken());
30                 context.write(word, one);
31             }
32         }
33     }
```

VS.

SELECT word,
COUNT(*) AS word_count
FROM words
GROUP BY word

SQL Don'ts in HIVE

SELECT * FROM ANYTHING: This brings back everything. Everything doesn't fit on a single computer.

JOIN: Join will take hours or days to perform and eat up all cluster bandwidth for everyone else trying to use it in the queue.

ORDER BY: Sorting is very computationally expensive.

Sub Queries: A sub query essentially creates a secondary table, which will be huge in HIVE.

Interactivity: SQL in DBMS is interactive because its almost instantaneous.

Execution Engine: Tez

The Stinger Initiative

2011, the world got together and declared MapReduce to be terrible.

- 44 companies
- 145 developers
- 392k lines of Java code



Hadoop 2.0 with Yarn & Tez

- Tez dropped Hive query times by **90%, 100x performance**
- Utilizes Apache Yarn
 - Yarn: resource manager for multi-cluster computing
- Introduced partial in-memory, local head nodes
- Rewrote HiveQL as an actual language, instead of translation

HADOOP IN THE AZURE CLOUD

Hadoop Implementations

hadoop




HDInsight



Amazon Elastic
MapReduce

data science dojo
— data science for everyone —

Hadoop in Azure




MACHINE LEARNING AT SCALE - REVISITED

What is Mahout?


- Distributed Machine Learning platform.
- Built on top of MapReduce and HDFS.
- Script-based and command line interfaces.
- R-like language implementation.



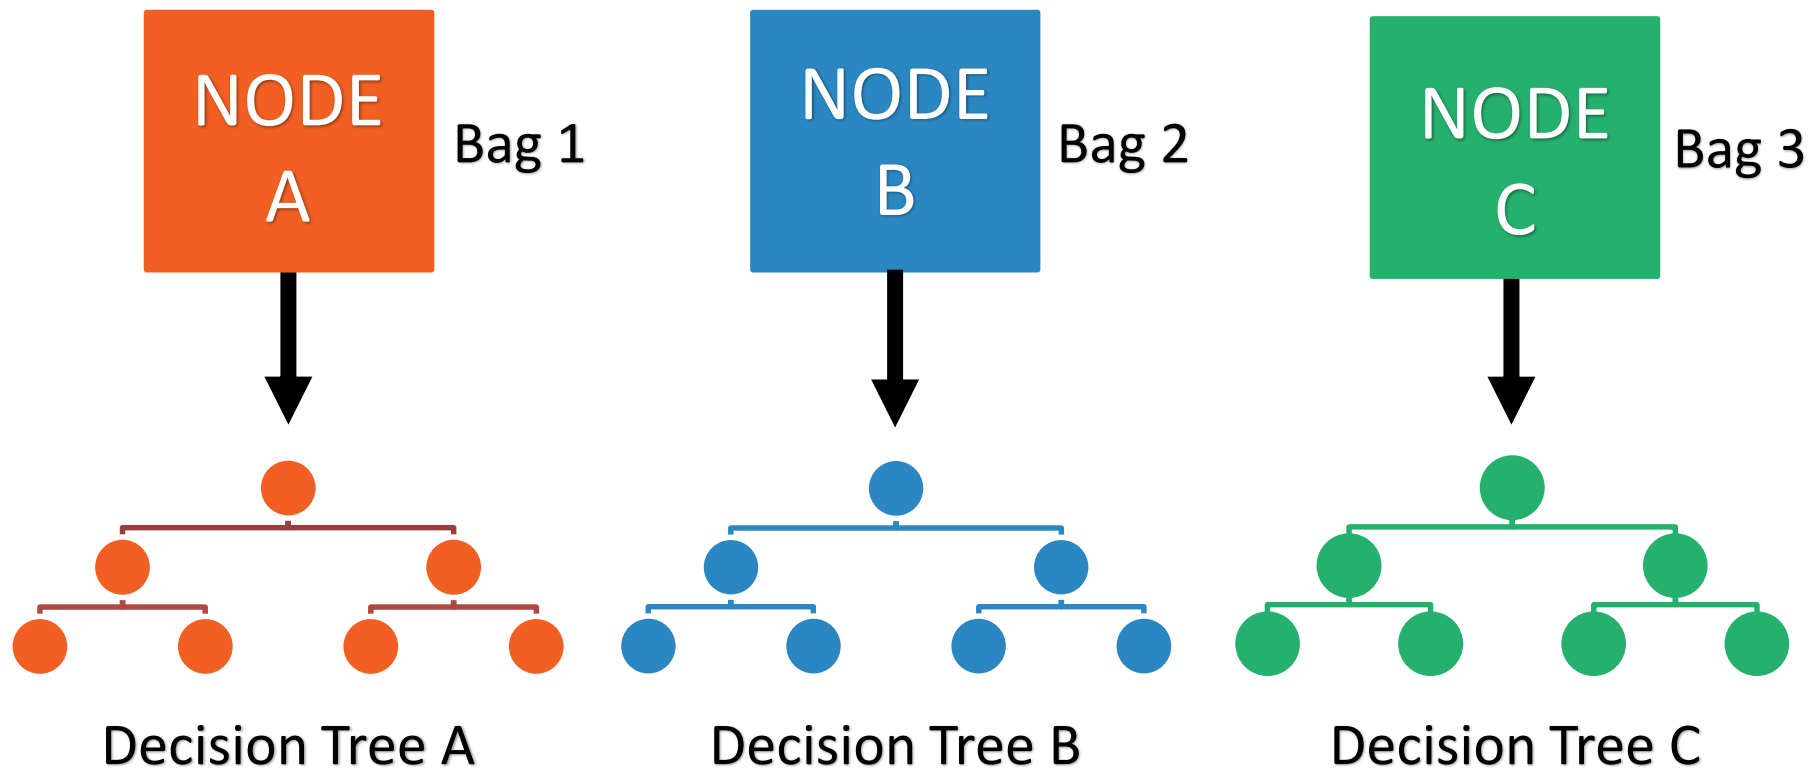
Distributed Random Forest



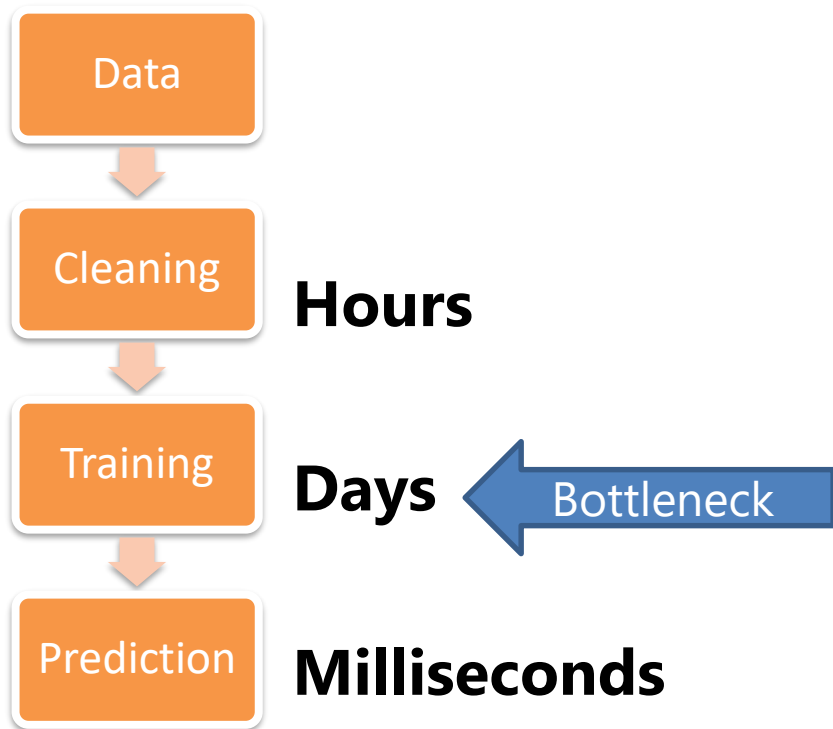
Distributed Random Forest



Distributed Random Forest



Processing Times - Machine Learning



- Large scale systems are only needed for training
- Phones can use models outputted by mahout to predict new data
- After a model is trained, save the model to any IO file type and reload it where you want

DISTRIBUTED COMPUTING V2.0 – APACHE SPARK


What is Spark?

- “A fast and general engine for large-scale data processing.”
- Designed to incorporate the goodness of Hadoop and address Hadoop’s shortcomings.
- Can complement Hadoop via integration with both HDFS and Hive.



Why Spark? Improved Perf!

- Up to 10x faster than Hadoop working with data from disk.
- Up to 100x faster working with data stored in memory!



Big Data, Faster!

3x faster on 10x fewer machines!

Daytona GraySort Contest: Sort 100 TB of data!

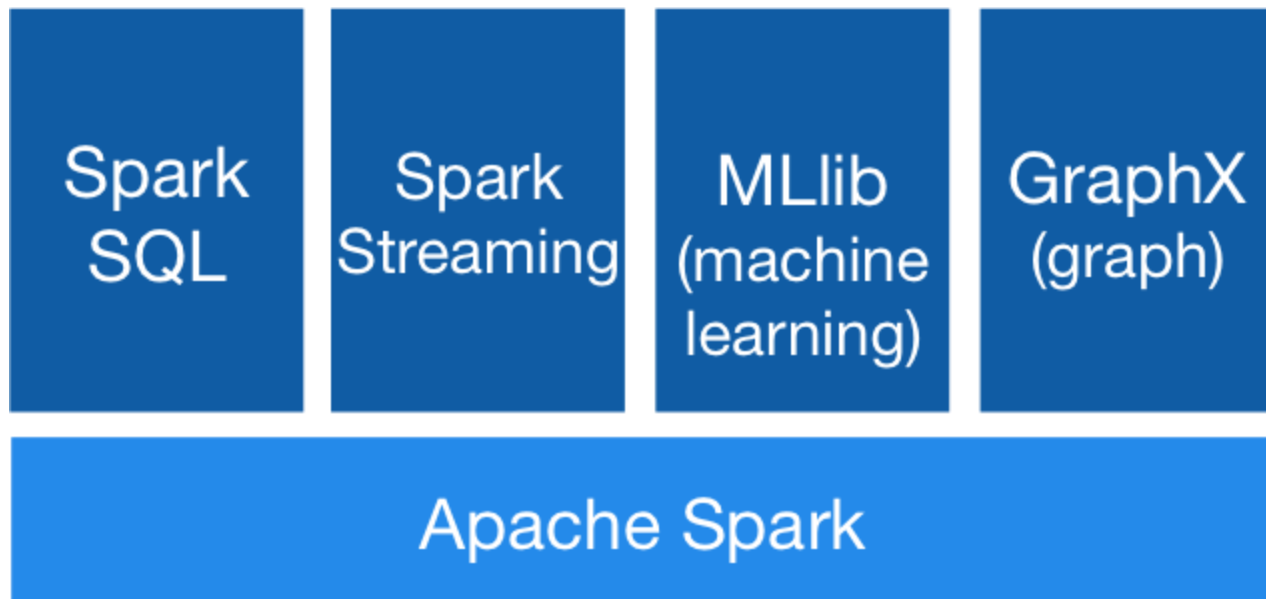
Previous World Record:

- Method: Hadoop
- Yahoo!
- 72 Minutes
- 2100 Nodes

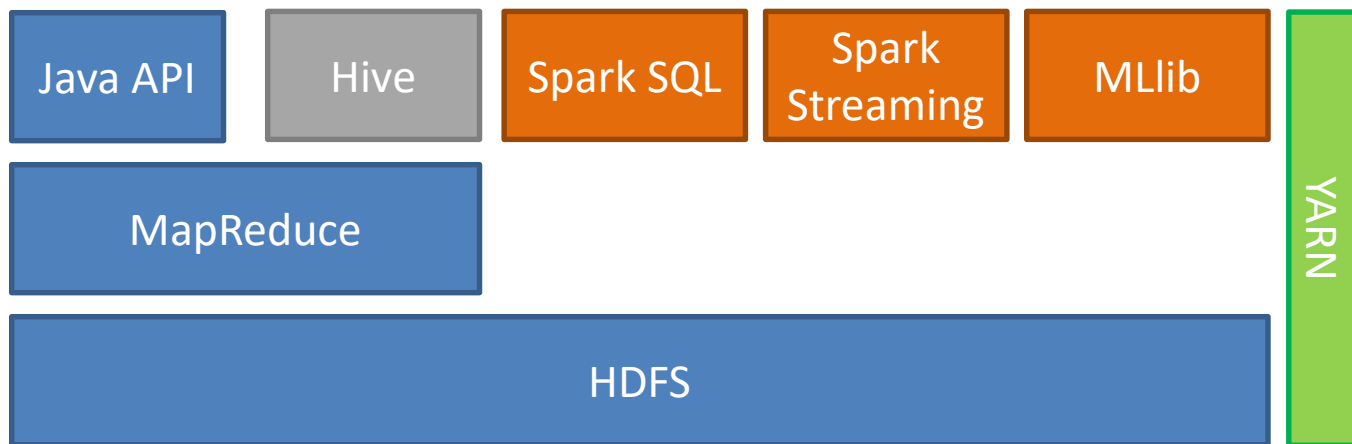
2014:

- Method: Spark
- Databricks
- 23 Minutes
- 206 Nodes

Conceptual Architecture



Spark and Hadoop




- Spark can be deployed on a Hadoop cluster and share cluster resources via YARN.
- Spark, however, does not require Hadoop!

Why is Spark Faster?

- First, Spark processing implements *lazy execution*:
 - Data operations are either *transformations* or *actions*.
 - Transformations are not executed immediately, but are stored.
 - When an action is issued, Spark evaluates all stored transformations and optimizes processing before executing.
- Second, Spark performs most processing in-memory:
 - RAM is far faster than using disk storage – even SSD drives.
 - More RAM in the cluster allows Spark to process data faster.

Technology adoption life cycle



Source: <http://carlosmartinezt.com/2010/06/technology-adoption-life-cycle/>

QUESTIONS

APPENDIX


MapReduce, via Playing Cards





Let's count the number of spades, clubs, hearts, and diamonds in a stack of cards, the way map reduce would.

- Each card represents a row of data
- Each suit & number represents an attribute of the data



Using a 2 Data Node Cluster




Mapping: Each Node's HDFS



Mapping: Node Sorting



Mapping: Node Shuffle, Data Transfer



Mapping: Node Shuffle, Data Transfer

